

Cloud On-Ramp (Zero → VM → GPU → K8s Job)

Requirements: cloud account (for VM/GPU), Docker+kind+kubectl, SSH key.

1. First VM (10–20 min)

- Create smallest Linux VM; open SSH (22) to your IP
- `ssh -i KEY.pem user@IP`
- Run hello Python script → Verify "hello, cloud"

2. First GPU Job (15–30 min)

- Create smallest GPU VM; `nvidia-smi`
- Run tiny PyTorch check → Verify CUDA True/False

3. First K8s Job (10–20 min)

- `kind create cluster --name onramp`
- `kubectl apply -f hello-job.yaml`
- `kubectl logs job/hello-job` → Verify message

Stop/terminate VMs when finished • Repo: